

Conditional Entropy Defenses with Attention

Slot + Gated Cross Attention vs. Gated Attention Pooling

Project Update

Group Meeting Preparation

Attention Privacy Project


Roadmap

- Recap: collaborative inference threat model and CEM requirements
- Slot + Gated Cross Attention surrogate
- Gated Attention Pooling surrogate
- Comparison, monitoring guidance, next steps

Background

Collaborative Inference & CEM Goal

- Local encoder F_e uploads feature vectors z to the cloud decoder F_d ; attacker can attempt input reconstruction.
- Conditional entropy maximization (CEM) seeks to keep $H(x | z)$ high so reconstructions remain uncertain.
- Legacy implementation used KMeans/GMM to approximate per-class variance and penalize sharp feature clusters.
- New attention surrogates replace GMM for richer class structure modeling while remaining differentiable.




Design Criteria for Attention CEM

- **Fine-grained:** capture multi-modal subclasses within each label instead of assuming a single Gaussian.
- **Stable:** avoid exploding gradients and NaNs; guard against small batches or class imbalance.
- **Optimizer-friendly:** expose gradients that can be merged with encoder updates via λ .
- **Diagnostics:** surface interpretable statistics (slot gates, SNR signals, variance margins) to guide tuning.

Slot + Gated Cross Attention

Pipeline Overview



- Processes one class at a time; reuses slots to encode latent sub-modes.
- Cross attention re-injects slot context into every sample before measuring dispersion.

Slot Attention Core (Locatello et al.)

- Inputs are LayerNormed and projected to keys/values once per class.
- Slots are initialized from a learned Gaussian (μ, σ) and iteratively refined ($T=3$).
- Competitive soft assignments $r_{ms} = \text{softmax}(\beta \text{ sim}(x_m, s_s))$ ensure each slot explains part of the class.
- GRU + MLP residual updates preserve temporal dynamics and stabilize slot representations.
- Learnable temperature β and slot sharpening power adjust how concentrated each slot becomes.

Flamingo-Style Gated Cross Attention

- Pre-LN multi-head cross attention (4 heads) with learnable α_{attn} gate:

$$y = q + \tanh(\alpha_{\text{attn}}) \cdot \text{CrossAttn}(q, s).$$

- Feed-forward block mirrors Flamingo dense layer with a second gate α_{ffn} for gradual activation.
- Slots act as shared KV memory; query tokens are the LayerNormed class samples.
- Dropout-free path keeps gradients deterministic, easing debugging.

Conditional Entropy Surrogate

1. **Mixture-of-slots variance:** compute per-slot mean μ_s and variance σ_s^2 under sharpened responsibilities.
2. **Per-dimension gate:** $\text{sigmoid}(\text{MLP}(\text{LN}(\log \sigma_s^2)))$ suppresses low-signal dimensions.
3. **SNR gate:** $\sigma(\kappa(\sigma_s^2/(\mu_s^2 + \epsilon) - \tau))$ dampens high-SNR activations that leak semantics.
4. **Softplus margin:** $L_{\text{base}} = \frac{1}{\beta'} \log(1 + e^{\beta'(\log \sigma_s^2 - \log \tau - m)})$ smooths the hinge at the variance threshold.
5. **Slot mass weighting:** responsibilities are sharpened by $w_s \propto (\text{mass}_s/M)^\rho$ to focus on dominant slots.
6. **Class gate:** $g_c = \sigma(a(M/B - b))$ scales classes with scant samples to avoid noisy penalties.

Training Integration & Safety Nets

- Activation: only after warmup (> 3 epochs), non-random centers, and $\lambda > 0$.
- Gradients: backprop through L_C , stash encoder gradients, zero optimizers, then blend with classification gradients using λ and LR-aware rescaling.
- Scaling: final L_C contribution multiplied by `attention_loss_scale = 0.25` before gradient capture.
- Early shutoff: for first 100 calls or when gate statistics exceed thresholds, L_C returns a graph-connected zero to avoid destabilizing the encoder.
- Robustness: skip computation if NaN/Inf features, and lazily register module parameters into the optimizer only once.
- Logging: prints mean gate activations and slot statistics every 200 calls to guide hyper-parameter tuning.

Gated Attention Pooling

Why a Second Surrogate?

- Slot pipeline is expressive but heavy: GRUs, multi-head attention, and numerous gates increase flops and tuning cost.
- Need a lightweight alternative for large batches, high-resolution encoders, or quick ablations.
- Gated Attention Pooling borrows from deep MIL: a single attention distribution summarizes class features.

Pooling Architecture



Mathematics

$$a_m = \frac{\exp(w^\top [\tanh(Vx_m) \odot \sigma(Ux_m)])}{\sum_j \exp(\cdot)},$$
$$\mu_c = \sum_m a_m x_m, \quad \sigma_c^2 = \sum_m a_m (x_m - \mu_c)^2,$$
$$\tau = \text{var_threshold} \cdot \text{reg_strength}^2 + \gamma,$$
$$L_C = \max\{0, \log(\sigma_c^2 + \gamma) - \log(\tau)\}.$$

- LayerNorm of inputs curbs variance collapse under sharp attention distributions.
- Hidden size defaults to $\min(512, \max(64, D/4))$ per class at instantiation.
- ReLU-style hinge keeps gradients sparse when class variance already exceeds the target threshold.

Training Notes

- Warmup extends to 5 epochs to let softmax weights stabilize.
- Module lazily constructed on first use and registered with the optimizer; gradients combined with encoder updates just like the slot version.
- Uses a smaller scaling factor: `attention_loss_scale = 0.1` before gradient capture.
- Skips classes with ≤ 1 samples to avoid ill-defined variance; safe defaults return zero loss.
- Same guardrails for NaN/Inf detection and optimizer resets after the auxiliary backward pass.

When to Prefer Each

- **Use Slot + Gated Cross Attention** when class heterogeneity is high (complex datasets, multi-modal privacy leakage) and logging bandwidth is available.
- **Use Gated Attention Pooling** for rapid experimentation, limited compute, or when you mainly need coarse variance inflation.
- Both surrogates can be toggled via the training script; only one runs per experiment.

Comparison & Next Steps

Side-by-Side Summary

	Slot + Gated Cross Attn	Gated Attention Pooling
Representation	Multi-slot latent mixture with cross-attention refinement	Single attention distribution over class tokens
Gates	Per-dim MLP, SNR gate, slot mass, class gate, softplus margin	Single gated attention + variance hinge
Compute	GRU iterations + multi-head attention (heavier)	Two linear layers + softmax (light)
Tuning knobs	β , slot power, SNR threshold, class gate bias, margin	Threshold τ , loss scale, hidden width
Safety	Early shutoff, extensive logging, NaN guards	Skip small classes, automatic hidden size

Operational Guidance

- Monitor: slot gate means, hard gate averages, and class-level activations printed by the module; large values signal over-sharpening.
- For Gated Pooling, track variance histograms versus τ to ensure the hinge is active for risky classes.
- Adjust λ schedule alongside `attention_loss_scale` to balance accuracy and privacy cost.
- Remember to rerun validation with reconstruction attacks after enabling either surrogate to quantify impact.

Thank you!